

Day 6: Industry Landscape: Cloud Providers & AI Chips

An exploration of the evolving infrastructure ecosystem powering today's AI revolution, from cloud service providers to the specialized silicon driving model training and inference.



The AI Infrastructure Ecosystem

The modern AI landscape is fundamentally shaped by two complementary forces:

Cloud Providers

Deliver on-demand compute, scalability, and managed AI services that democratize access to AI capabilities

Chip Manufacturers

Supply the raw computational horsepower required for increasingly complex model training and inference workloads

This competitive ecosystem drives rapid innovation cycles and improving cost efficiency — making strategic infrastructure knowledge essential for AI teams.



Cloud Providers Leading AI Infrastructure

AWS

- EC2 P4/P5 GPU instances
- SageMaker suite for ML workflows
- Largest market share in cloud

Google Cloud

- Custom TPU accelerators
- Vertex AI end-to-end platform
- AI-optimized networking infrastructure


Microsoft Azure

- NVIDIA H100 GPU clusters
- Exclusive OpenAI partnership
- Enterprise-focused AI solutions

Emerging Specialists

A new breed of [specialized AI cloud providers](#) like CoreWeave and Lambda Labs is disrupting the market with GPU-first architectures and optimized AI infrastructure, challenging the traditional hyperscalers.

AI Chip Makers & Hardware Vendors



NVIDIA

Dominant market leader with H100 GPUs commanding premium pricing; de facto standard for AI training



Google TPUs

Custom tensor processing units optimized for TensorFlow and large language models



AMD

Instinct MI300 series challenging NVIDIA with competitive performance at lower power consumption



Startups & Specialists

Cerebras (wafer-scale), Graphcore (IPUs), SambaNova (dataflow architecture) pursuing novel approaches

Cloud + Chip Synergy

The relationship between cloud providers and chip manufacturers is increasingly **symbiotic and strategic**, shaping where cutting-edge AI development happens.

AWS + NVIDIA

Purpose-built EC2 instances featuring massive H100 GPU clusters with optimized networking fabric

Google Cloud TPU Exclusivity

TPU v4 pods available only on GCP, creating a compelling ecosystem for TensorFlow users

Azure + OpenAI

Exclusive access to OpenAI's infrastructure and models drives enterprise adoption of Azure



What It Means for AI Engineers

Success in the modern AI landscape requires strategic infrastructure decisions and adaptable skills:

Multi-Environment Expertise

Engineers must develop proficiency across cloud platforms and accelerator architectures

Balance Lock-in vs. Performance

Weigh vendor-specific optimizations against flexibility and portability concerns

Cost-Performance Optimization

Continuously evaluate infrastructure choices as chip performance doubles roughly every 2 years

Engineers who understand this landscape gain a [significant career advantage](#) in architecting next-generation AI systems.

