

Scaling Storage for AI Training Datasets

Essential strategies for ML engineers to overcome storage bottlenecks and optimize AI training infrastructure




Why Storage Scaling Matters for AI

AI training requires **terabytes to petabytes** of data, with storage systems that must keep pace with both **data growth & training speed**.

Poor storage architecture leads to:

- Slower training epochs
- Wasted GPU computation hours
- Higher operational costs

Finding the right balance between **cost, performance, and scalability** is critical for LLMs, computer vision, and generative AI training pipelines.



Storage Bottlenecks in AI Training

Slow I/O Throughput

GPUs sit idle waiting for data, wasting expensive compute resources and extending training time

Network Limitations

Bandwidth constraints delay data delivery to training nodes, creating pipeline stalls


Dataset Duplication

Multiple copies waste storage capacity and increase infrastructure costs unnecessarily

Inefficient Caching

Poor caching strategies cause repeated downloads from slower storage tiers

When storage can't keep up, training workloads stall, creating a costly domino effect across your AI infrastructure.



Storage Scaling Strategies



Scale Up

- Faster SSDs/NVMe
- Optimized filesystems
- Vertical growth



Scale Out

- Distributed storage
- Horizontal expansion
- Multi-node access



Cloud Object Storage

- Near-infinite capacity
- Elasticity
- Pay-as-you-go



Caching + Sharding

- Reduced I/O strain
- Parallel access
- Locality benefits

For production AI environments, a **hybrid approach** combining these strategies delivers the most balanced performance and cost efficiency.

Distributed File Systems for AI Workloads

1

HDFS (Hadoop Distributed FS)

Mature ecosystem with big data roots. Handles large datasets but not optimized for GPU workloads. Good for preprocessing pipelines.

2

CephFS / GlusterFS


Open-source scale-out systems with flexibility. Software-defined storage that can run on commodity hardware. Popular in enterprise AI clusters.

3

Lustre

HPC-optimized parallel filesystem with excellent throughput for large files. Complex deployment but delivers exceptional performance for distributed training.

These systems provide a [shared namespace across nodes](#), enabling unified dataset access from any compute node in your AI cluster.




Cloud Object Storage for AI

Key Players

- AWS S3
- Google Cloud Storage (GCS)
- Azure Blob Storage

Benefits for AI Workloads

- Virtually [infinite capacity scaling](#)
- Cost-effective for raw dataset storage
- API-driven access from training jobs
- Integrated data versioning and lifecycle policies



Ideal Use Cases:

- Dataset lakes for long-term archiving
- Staging area before preprocessing
- Shared repositories across research teams

High-Performance Storage for Training

1

NVMe SSDs

Deploy on local training nodes for hot data access. 20-30x faster than HDDs with minimal latency. Essential for LLM and vision model fine-tuning.

2

Parallel File Systems

Lustre and BeeGFS deliver distributed access with high throughput. Used by top AI research labs for multi-node training.

3

Cloud Premium Tiers

S3 Express One Zone, GCS Turbo offer millisecond-level latency for critical workloads. Higher cost but better performance.

Data locality = faster access for GPU clusters

These high-performance solutions are particularly crucial for large language models and vision workloads where data throughput directly impacts training efficiency.



Caching and Preprocessing Strategies

Techniques for Optimizing Data Flow

- **Local Caching:** Keep frequently used datasets on fast node-local storage
- **Pre-sharding:** Split large datasets for parallel reads across nodes
- **Format Optimization:** Use binary formats optimized for ML workloads

Key Tooling

- [WebDataset](#): Efficient loading of tar files for PyTorch
- [TFRecords](#): Optimized binary format for TensorFlow
- [Hugging Face Datasets](#): Memory-mapped access for transformers



Proper caching reduces pressure on backend storage and dramatically improves GPU utilization efficiency.

Scaling Best Practices

Implement Hybrid Tiered Storage

Combine hot SSD tiers for active training with cold object storage for archival. Automatically migrate data between tiers based on access patterns.

Monitor I/O Throughput

Use Prometheus + Grafana dashboards to track storage performance metrics. Set alerts for bottlenecks before they impact training jobs.

Deduplicate Datasets

Implement content-based deduplication to eliminate redundant data. Can reduce storage requirements by 20-30% for computer vision datasets.


Optimize Dataset Formats

Structure data for sequential reads to match training batch patterns. Use columnar formats for tabular data and binary formats for images/embeddings.

Automate with Orchestration

Use data orchestration frameworks like Airflow or Prefect to manage preprocessing and data movement between storage tiers.

Real-World Examples



OpenAI GPT Training

Uses sharded dataset pipeline from object storage with custom caching layer to feed thousands of GPUs simultaneously

Google DeepMind

Deploys HPC storage clusters for reinforcement learning simulations with ultra-low latency requirements

Tesla Autopilot

Manages petabyte-scale vision datasets on distributed file systems with specialized video preprocessing pipelines

Netflix Recommendations

Combines hybrid file system and cloud storage architecture for their recommendation model training pipelines

The industry consistently relies on **multi-layered storage systems** that balance performance, scalability, and cost-efficiency.

Key Takeaways



AI training demands scalable, high-performance storage

Storage throughput directly impacts GPU utilization and training time



Combine distributed FS + cloud object storage

Hybrid approach offers best balance of flexibility, performance and cost



Implement caching + data sharding

Properly optimized data pipelines can improve I/O efficiency by 5-10x

Your storage strategy is as critical as your GPU architecture in building efficient AI infrastructure. Optimize now to accelerate your ML development cycle and reduce operational costs.

