




Multi-GPU Scaling and Interconnects (NVLink)

An essential guide for ML engineers and infrastructure architects building advanced multi-GPU training systems for large-scale AI workloads.

Why Multi-GPU Scaling?

Modern AI development faces fundamental constraints that single GPU solutions simply cannot overcome:

Breaking Physical Limits


Single GPU memory and compute ceilings constrain model size and training capability

Enabling Scale

Large AI models (LLMs, computer vision, reinforcement learning) demand distributed computation


Accelerating Development

Multi-GPU configurations deliver faster training cycles and support larger batch sizes



Multi-GPU systems are the foundation of trillion-parameter model training

Data Parallelism: The Foundation of Multi-GPU Training



Core Principles

- Mini-batches are **split across multiple GPUs**
- Each GPU has a complete model copy
- Gradients computed independently on each device
- Periodic synchronization combines results

□ Data parallelism is the most widely implemented technique in PyTorch and TensorFlow training pipelines, forming the cornerstone of distributed training.

Performance critically depends on **high-bandwidth GPU-GPU interconnects** for gradient synchronization.

Model Parallelism: Beyond Memory Constraints



Model > Single GPU Memory

When parameters exceed available VRAM, model must be partitioned



Layer Distribution

Different model components distributed across multiple devices




Coordinated Execution


Each GPU handles specific portions of forward/backward pass

Implementation Approaches

- **Pipeline Parallelism:** GPipe, PipeDream, DeepSpeed
- **Tensor Parallelism:** Splitting individual layers
- **Sequence Parallelism:** Distributing sequence dimensions

Model parallelism introduces significant **communication complexity** and requires careful orchestration of operations.





PCIe

NVLink

Interconnect Comparison: PCIe vs NVLink

PCIe (Peripheral Component Interconnect Express)

- Industry standard interconnect for CPU ↔ GPU communication
- Bandwidth ceiling of ~64 GB/s (PCIe 5.0 x16)
- Higher latency, limited peer-to-peer capabilities
- Becomes bottleneck in multi-GPU training workloads

NVLink (NVIDIA's High-Speed Interconnect)

- Purpose-built for GPU-to-GPU communication
- Bandwidth up to 900 GB/s with H100 GPUs
- ~14x faster than PCIe in peak scenarios
- Dramatically reduces synchronization overhead

NVLink Architecture

Direct GPU-to-GPU Communication


NVLink creates high-bandwidth, low-latency point-to-point connections between GPUs, enabling:

- Direct memory access between GPUs ([peer-to-peer](#))
- Elimination of CPU as communication middleman
- Significant reduction in data transfer overhead
- Support for Unified Memory Addressing (UMA)

NVLink's direct connections unlock **near-linear scaling** for communication-intensive workloads like transformer training.



Each NVLink connection provides dedicated bandwidth between GPU pairs



NVSwitch: The Multi-GPU Fabric



Extending NVLink's Reach

Functions like a network switch for GPUs, enabling scaling beyond direct connections



All-to-All Communication

Enables any GPU to communicate with any other GPU at full NVLink bandwidth




AI Supercomputing

Powers NVIDIA DGX, HGX systems and forms the backbone of modern AI supercomputers

NVSwitch removes communication bottlenecks in large GPU clusters, allowing **hundreds of GPUs** to function as a single computational unit.

Multi-GPU Topologies



Single Node, Multi-GPU

Up to 8×A100/H100 GPUs in systems like NVIDIA DGX, connected via NVLink/NVSwitch

Multi-Node Clusters

GPUs distributed across multiple servers, connected via high-speed networking (InfiniBand, RoCE)

Optimization Layer

NCCL (NVIDIA Collective Communications Library) optimizes operations for specific topologies

The topology selection directly impacts training throughput, scalability, and hardware utilization efficiency.

Bottlenecks in Multi-GPU Training

Communication Overhead

Gradient synchronization becomes increasingly expensive as GPU count grows, potentially negating computational benefits

Latency Limitations

When communication latency exceeds computation speed, GPUs spend time waiting rather than computing

Workload Imbalance


Uneven distribution of computation across GPUs creates stragglers that delay synchronization steps

Interconnect Saturation

PCIe-only configurations severely underutilize GPU computational capacity, creating bandwidth bottlenecks

Effective multi-GPU scaling requires **high-bandwidth, low-latency interconnects** to minimize these bottlenecks

Multi-GPU GPU Scaling Efficiency



Best Practices for Multi-GPU Scaling



Prioritize NVLink/NVSwitch-enabled hardware when budget permits

Optimize batch size to balance computation and communication

Implement mixed precision (FP16/BF16) to reduce communication volume

Leverage NCCL for topology-aware collective operations

Benchmark at small scale before deploying to hundreds of GPUs

i For large clusters, gradient accumulation can help reduce synchronization frequency while maintaining effective batch size.

Key Takeaways

Necessity of Scale

Multi-GPU configurations are no longer optional for large-scale AI training—they're essential infrastructure

Interconnect Hierarchy

PCIe provides baseline connectivity, while NVLink enables high-performance scaling with up to 14x higher bandwidth

Network Fabric

NVSwitch technology extends NVLink capabilities to create GPU superclusters with full-bandwidth connectivity

System Architecture

The combination of physical topology and communication libraries determines overall training efficiency

Understanding GPU interconnect technologies is the foundation of **mastering multi-GPU infrastructure**

Next steps: Evaluate your workloads, benchmark on different configurations, and design your infrastructure to balance performance and cost.