



Challenges of Multimodal Model Serving


Scaling AI Across Text, Images, Audio, and Video

A technical deep-dive for infrastructure engineers and ML platform teams responsible for deploying complex AI systems at scale.

Why Multimodal Serving Is Hard

Building infrastructure for single-modality models is already challenging. Multimodal systems compound these challenges by requiring:

- Processing of **different input formats** simultaneously
- Handling varied workloads from lightweight text to heavyweight video
- Meeting strict latency expectations for interactive applications



Your serving infrastructure must unify these diverse modalities while maintaining performance and reliability.

Input Complexity



Text

Tokenization pipelines, embedding generation, and vocabulary management



Images

Resizing, augmentation, CNN/ViT encoding with different resolution requirements



Audio

Spectrogram generation, ASR preprocessing, and feature extraction



Video


Frame extraction, temporal alignment, and motion analysis

Each modality requires specialized preprocessing, and [synchronization across these pipelines](#) adds another layer of complexity.

Model Fusion Challenges

Multimodal systems employ various fusion architectures:

- **Dual-encoder models** (like CLIP) process modalities separately then combine
- **Cross-attention fusion** allows modalities to influence each other's processing
- These approaches create significant memory and compute demands
- System designers must balance throughput with accuracy requirements




Infrastructure Constraints




Specialized GPUs

Different hardware optimizations for NLP vs. computer vision vs. speech recognition



Decoding Bottlenecks

Video/audio decoding is computationally expensive compared to text tokenization




Batching Complexity

Efficiently batching heterogeneous requests across modalities is non-trivial

Multi-tenant environments further complicate resource allocation with [unpredictable and bursty workloads](#).

Latency & Throughput Issues



The latency gap between modalities creates significant challenges:

- Text-only queries: typically **milliseconds**
- Video queries: often **seconds** without optimization

Cross-modal requests compound these issues:

- Bottlenecks are amplified when modalities must wait for each other
- Meeting p95/p99 latency SLAs becomes increasingly difficult

Scaling & Deployment Challenges



Memory Constraints

Deploying multiple specialized encoders significantly increases GPU memory requirements



Orchestration Complexity

Deciding which GPU handles which modality requires sophisticated routing logic



Varied Autoscaling Needs


Different modalities have distinct scaling patterns and resource requirements



Microservice Architecture

Need for containerized, modality-specific services that can scale independently

Monitoring & Observability



Diverse metrics across modalities complicate monitoring:

- Speech: WER (Word Error Rate)
- Text: BLEU, perplexity
- Images: Recall@k, precision

Additional challenges include:

- **Drift monitoring** across varied inputs (audio noise, image lighting)
- End-to-end tracing for multi-hop, cross-modal queries
- Unifying quality metrics into coherent dashboards

Security & Governance

Privacy Concerns

Video/audio inputs frequently contain PII/PHI that must be protected

Requires modality-specific privacy filters (face blurring, voice anonymization)

Compliance Requirements

GDPR, HIPAA, and other regulations apply differently across media types

Different data retention and processing requirements per modality

Audit Challenges

Logging policies must account for all modalities while balancing storage costs

Explainability becomes more complex with interacting modalities

Security strategies must be [comprehensive yet modality-aware](#) to properly protect sensitive data.

Best Practices

Modular Architecture

Build separate services for text, image, audio processing with clean APIs

Strategic Caching


Cache embeddings and frequent queries per modality to reduce computation

Model Optimization

Apply compression techniques for multimodal fusion models
(quantization, pruning)

Regional Deployment

Regionalize infrastructure to reduce bandwidth costs for media-heavy workloads



Key Takeaways



Unique Challenges

Multimodal serving faces distinct issues with input diversity, latency variance, and complex scaling requirements



Infrastructure Design

Systems must orchestrate specialized encoders and fusion components efficiently across heterogeneous hardware



Operational Complexity

Monitoring and governance are significantly more complex than with single-modality systems

Success requires a **thoughtful architecture** with modular services, strategic caching, and comprehensive compliance controls.

Contact us: platform-team@company.com